

Devaansh Sinha

(608)-335-0973 | devaansh.tech | sinhadevaansh@gmail.com | linkedin.com/in/devaanshsinha | github.com/devaanshsinha

EDUCATION

University of Wisconsin Madison

Sep. 2022 – Dec. 2025

Bachelor of Science in Computer Science, Minor in Game Design

GPA: 3.92 / 4.0

- *Relevant Coursework:* Operating Systems, Computer Vision, Big Data, Algorithms, Data Structures, Machine Organization, Intro to AI, Object Oriented Programming, UI/UX

TECHNICAL SKILLS

Languages: Java, Python, TypeScript, JavaScript, SQL, C/C++, HTML/CSS, Swift, C#

Frameworks & Libraries: Spring Boot, Django, gRPC, React, Next.js, REST APIs, Node.js, PyTorch, Tensorflow

Technologies & Tools: AWS, GCP, Docker, Kubernetes, Kafka, Spark, Hadoop, Terraform, CI/CD, PostgreSQL, MySQL, Git, JIRA, GitHub, Linux/Unix, Agile, Scrum, Bazel, SDLC, Unit Testing, LLMs, Claude code, Jenkins, Distributed Systems, OpenAI API, Embeddings models

EXPERIENCE

Gemini

New York City, NY

Software Developer

Feb 2026 – Present

- Worked on internal full-stack admin platform using TypeScript, React, and Scala, strengthening error handling and improving system observability
- Added and refined gRPC service endpoints to surface operational metadata and streamline internal access to system state

Software Engineer Intern

Jun 2025 - Feb 2026

- Redesigned Django-based backend microservice handling token activity, improving monitoring efficiency by 70% within a production system
- Stabilized Kubernetes workloads in a production environment by tuning CPU/memory requests and limits, eliminating OOMKills and reducing pod restart rates
- Designed automated token pruning pipeline reducing PostgreSQL storage by 200GB and improving API throughput by 40%
- Led upgrade of Stripe payment APIs, refactoring legacy endpoints and implementing new billing workflows to improve transaction reliability and maintain compliance with updated SDK standards

Software Engineer Intern

Jun 2024 - Aug 2024

- Built internal automation platform used by 800+ engineers, eliminating 10+ weekly hours of manual reporting via Django, AWS, and PostgreSQL, built for SDLC and Agile methodologies
- Developed an interactive team management dashboard within Slack with configurable reminders, automated prompts, and visualized work summaries, improving engineering lead oversight and team transparency
- Owned backend features end-to-end across technical design, implementation, testing, code reviews, and production releases within an Agile SDLC environment

Nandi Travels

Niles, IL

Full Stack Developer

May 2025 – Present

- Designed and developed production-grade full-stack platform using React, Node.js, and PostgreSQL, exposing RESTful microservices supporting booking workflows and real-time updates for 1K+ monthly users
- Built event driven automation using AWS Lambda and Twilio to handle 200+ weekly booking and trip notifications, reducing manual operations and improving system reliability

Bajaj Finserv

Pune, India

Software Engineer

Jun 2023 – Aug 2023

- Developed a spam detection microservice using Spring Boot, safeguarding 500K+ loan applications with 95% classification accuracy.
- Developed highly scalable Spring Boot microservice safeguarding 500K+ loan applications, contributing to mission-critical backend systems

PROJECTS

Remnis

- Built a privacy first macOS desktop app using Tauri and a Python sidecar that captures developer workflow locally to create a searchable history, enabling users to query past tasks, errors, and work context
- Implemented a local semantic search pipeline using MiniLM embeddings and LanceDB, enabling time-aware and intent-based queries over a user's workflow history
- Added an on demand AI layer combining vector retrieval, reranking, and a small local LLM to generate grounded summaries from retrieved evidence without sending data off device

Cogniflow

- Designed distributed backend data pipelines aggregating blockchain transactions into a scalable knowledge graph, enabling real-time analytics and API-driven querying
- Built vector-based retrieval pipelines using Supabase pgvector and OpenAI embeddings to enable semantic search, wallet clustering, and contract behavior detection across EVM chains
- Implemented event-driven indexers with TypeScript workers to continuously sync on-chain data, compute token-level activity metrics, and surface insights through an AI-powered query layer

Shelfwise

- Built comprehensive food tracking application using Python/OpenCV for receipt scanning with 98% OCR accuracy
- Developed scalable FastAPI microservice backend with Prisma ORM supporting real-time inventory synchronization and API-driven workflows and integrated OpenAI powered NLP assistant for hands free interaction and intelligent recipe suggestions